Serverless Data Pipeline

Data Pipeline :

The purpose of this practical is to design and implement a data pipeline on a server that can collect data from different sources, process it, transform it, and finally load it into a storage system such as a database, data warehouse, or cloud storage. This involves using server-side scripting, automation, logging, and error-handling techniques.

Key Steps of the Practical -

- Store orders in **DynamoDB**
- Use **Glue** to crawl + transform data
- Store cleaned data in **S3 (Data Lake)**
- Use **Athena** to run SQL queries on the data

Steps for Practical

Step 1 : Create DynamoDB Table

- Open DynamoDB → Create Table
- Table name: Orders
- Partition Key: OrderID (String)
- Sort Key: OrderDate (String) → Create Table

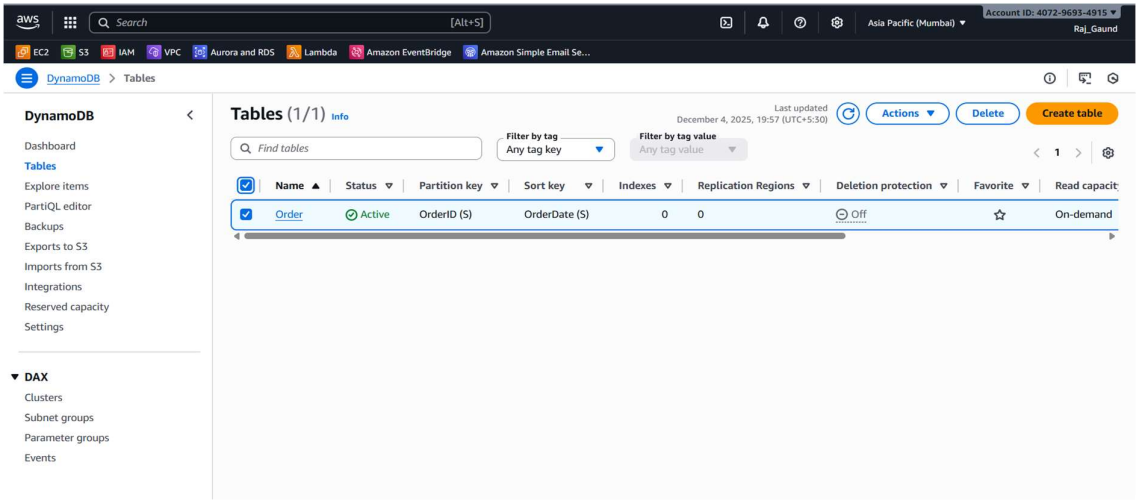

A new DynamoDB table is created to store order records.

Step 2: Insert Sample Data

- Click on Table (Order) → Explore Item
- Create Item → Add Attribute and Data → Create
- Add at Least five items

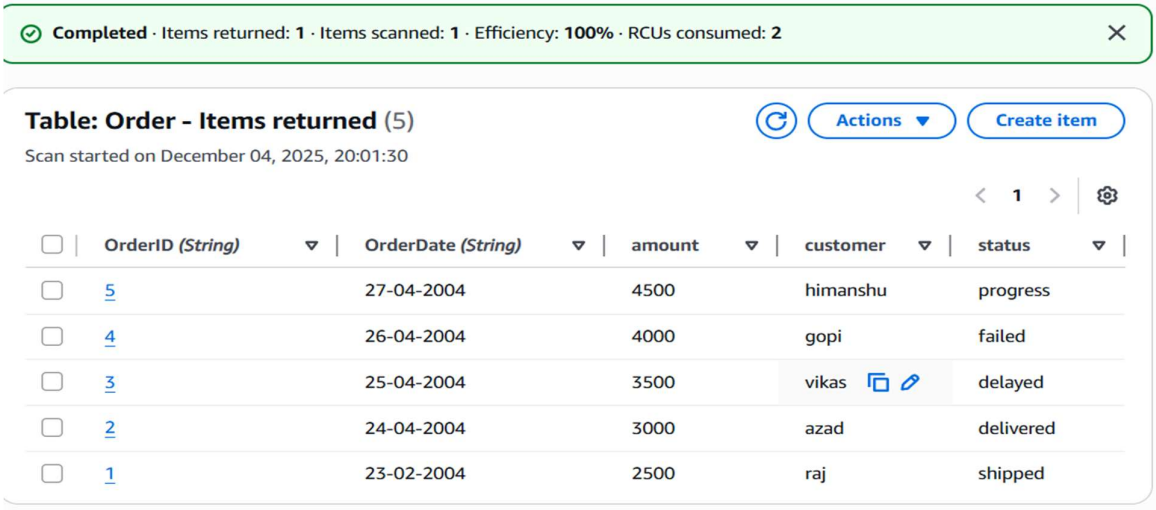

DynamoDB now contains real order data that the pipeline will process.

Step 3: Create Glue Crawler

- AWS Console → **Glue** → Data Catalouge
- Create Crawler → Enter Crawler Name (OrderCrawler)
- Add Datasource → Data source : **DynamoDB**
- Table name : DynamoDB Table Name (**Order**) → Add
- Choose IAM Role : Create New IAM role
- AWS → IAM → Role → Create → Name : GlueCrawlerRole
- Trusted Entity : AWS service → Use case : **Glue**
- Add Policies / Permission : **AWSGlueServiceRole** , **AmazonDynamoDBReadOnlyAccess , AmazonS3FullAccess**
- Create IAM Role
- Add Database : **ordersdb** → OnDemand → Create Crawler

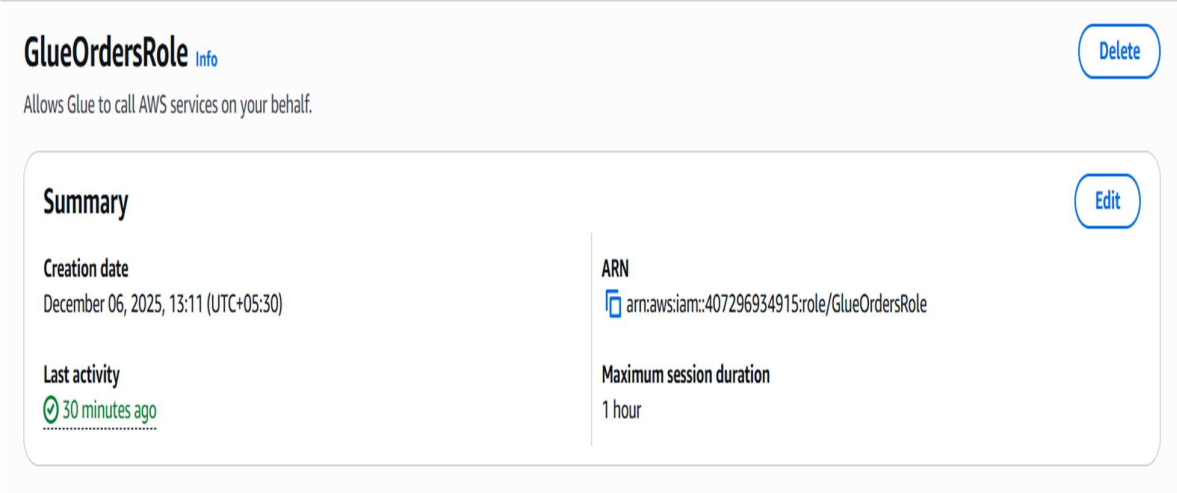

IAM Role Created Successfully

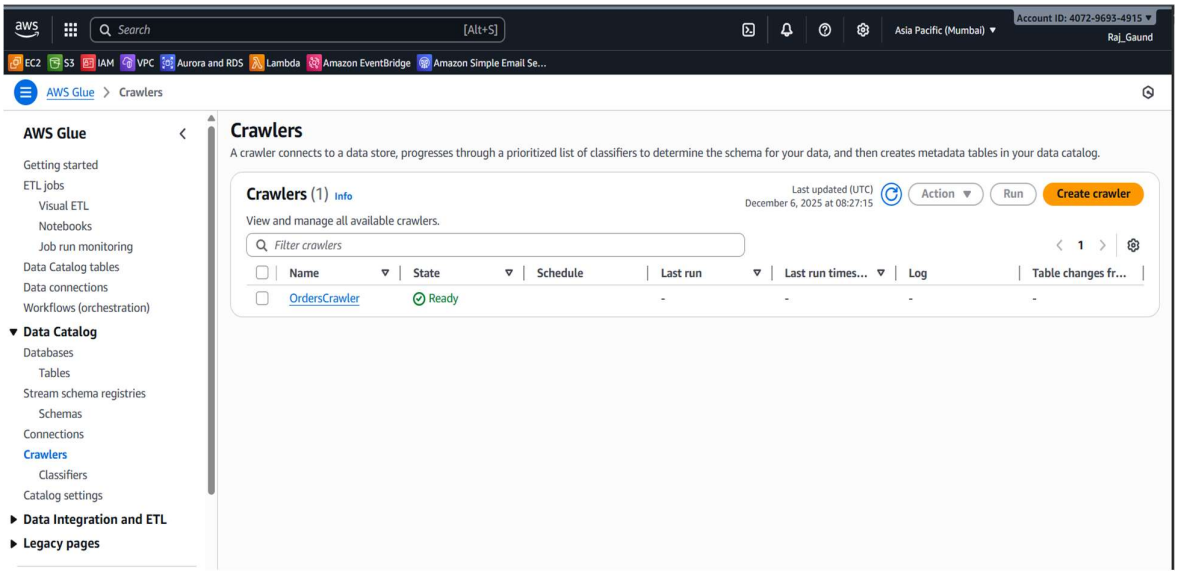

Crawler create Successfully

Step 4 : Run The Crawler

- AWS → Glue → Crawler
- Click Crawler → **Run Crawler**
- AWS → Glue → Table

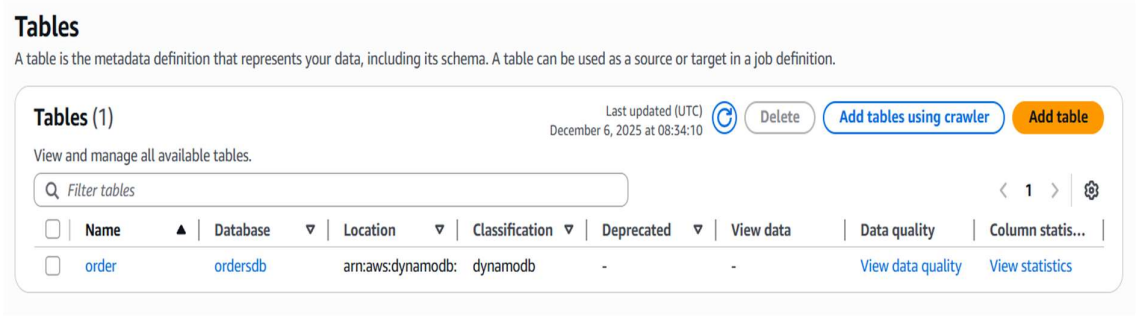

(Table created automatic when you run the crawler)

Step 5 : Create S3 bucket (for ETL output)

- AWS Console → **S3** → General Purpose Bucket
- Bucket Name : you can enter own (my-orders-bucket123)
- Create Bucket

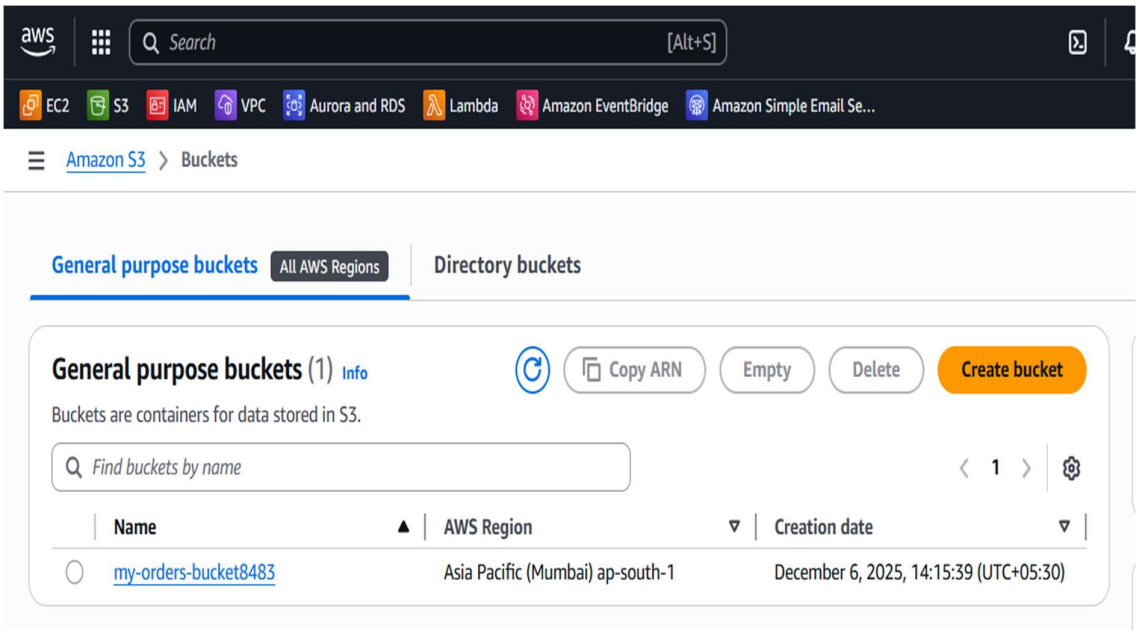

Bucket used to stored Job ETL output

Step 6 : Glue ETL Job create

- AWS → Glue → ETL jobs → Script editor
- Engine : Spark → Create script
- Type this code

Code :

```
import sys
from awsglue.transforms import Filter
from awsglue.utils import getResolvedOptions
from pyspark.context import SparkContext
from awsglue.context import GlueContext
from awsglue.job import Job

# required argument JOB_NAME
args = getResolvedOptions(sys.argv, ['JOB_NAME'])
sc = SparkContext()
glueContext = GlueContext(sc)
spark = glueContext.spark_session
job = Job(glueContext)
job.init(args['JOB_NAME'], args)

# ----- CONFIG: change these if you used other names -----
GLUE_DATABASE = "ordersdb"          # Glue DB created by crawler
GLUE_TABLE = "order"                # Glue table name (crawler output)
S3_OUTPUT_PATH = "s3://my-orders-bucket123/shipped/"  # change to your bucket
```

```
datasource = glueContext.create_dynamic_frame.from_catalog(
    database=GLUE_DATABASE,
    table_name=GLUE_TABLE
)

# Optional: print schema to logs (helpful)
print("Schema: ", datasource.schema())

# Filter function: keep only Status == 'Shipped'
def is_shipped(rec):
    status = rec.get('Status') or rec.get('status') or rec.get('STATUS')
    if status is None:
        return False
    try:
        return str(status).lower() == 'shipped'
    except:
        return False

filtered = Filter.apply(frame=datasource, f=is_shipped)

# Write filtered data to S3 as Parquet
glueContext.write_dynamic_frame.from_options(
    frame=filtered,
    connection_type="s3",
    connection_options={"path": S3_OUTPUT_PATH},
    format="parquet"
)

job.commit()
```

- Jobs details → Enter name : OrdersETLJob
- IAM Role : Choose existing (GlueOrdersRole)
- Glue version: Glue 3.0 (default) / Engine = Spark (keep)
- Language : Python → Save → Run Job
- Click on OrdersETLJob → job details → Run

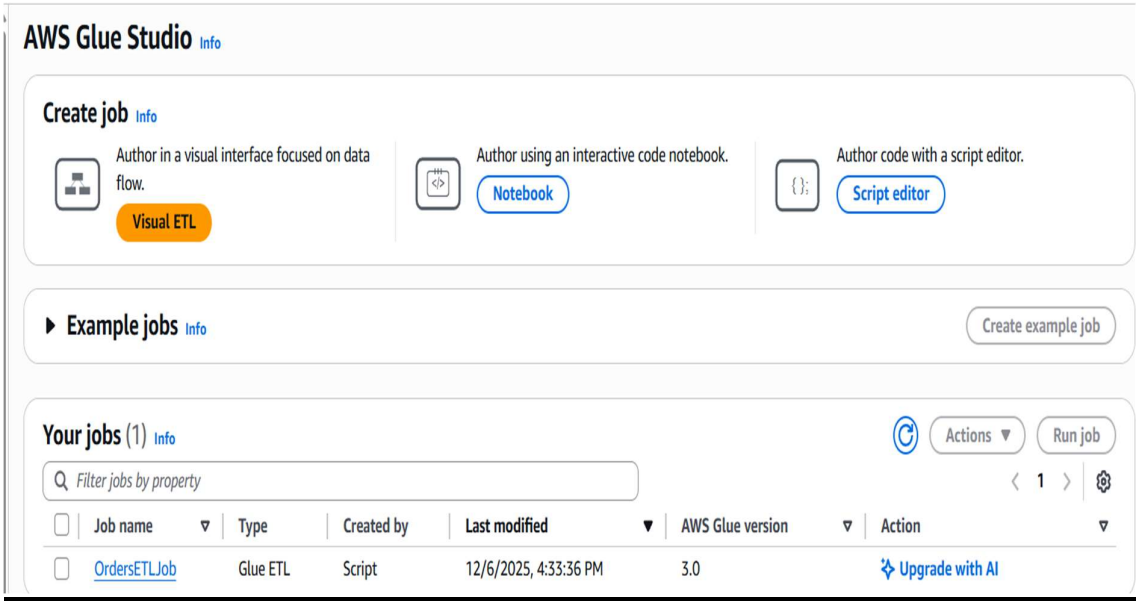

If succeed show means

- Script is correct
- DynamoDB read that Data
- Filter is successfully applied

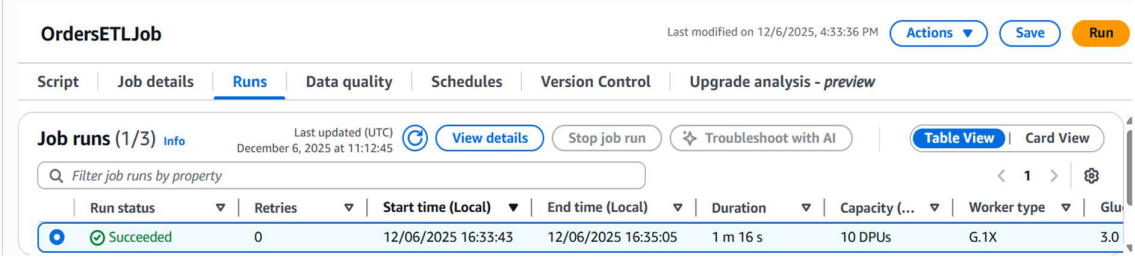

Step 7 : Open S3 bucket

- AWS → S3 → Open Bucket
- Example : (my-orders-bucket8483)

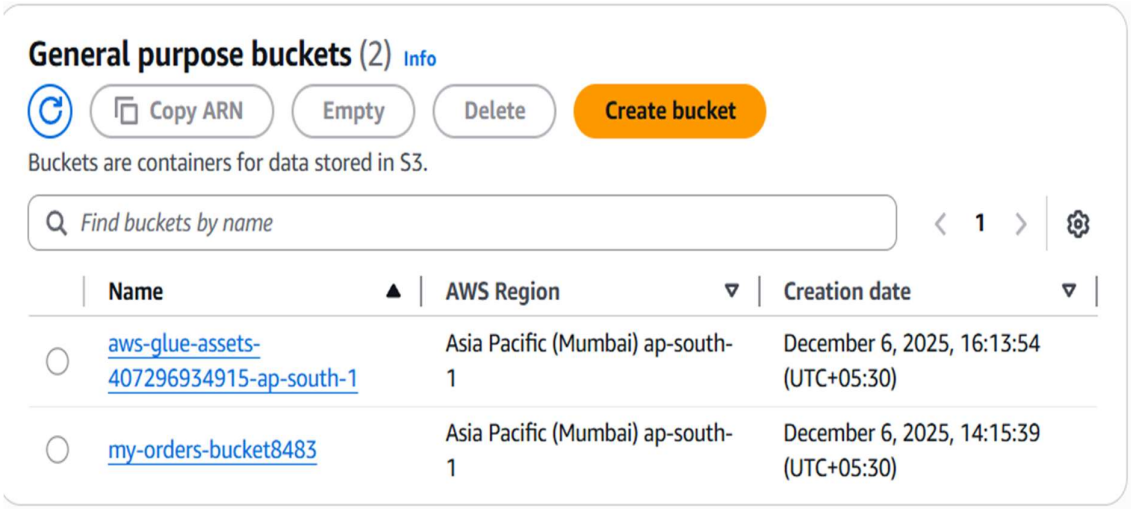

- Clicked shipped named folder

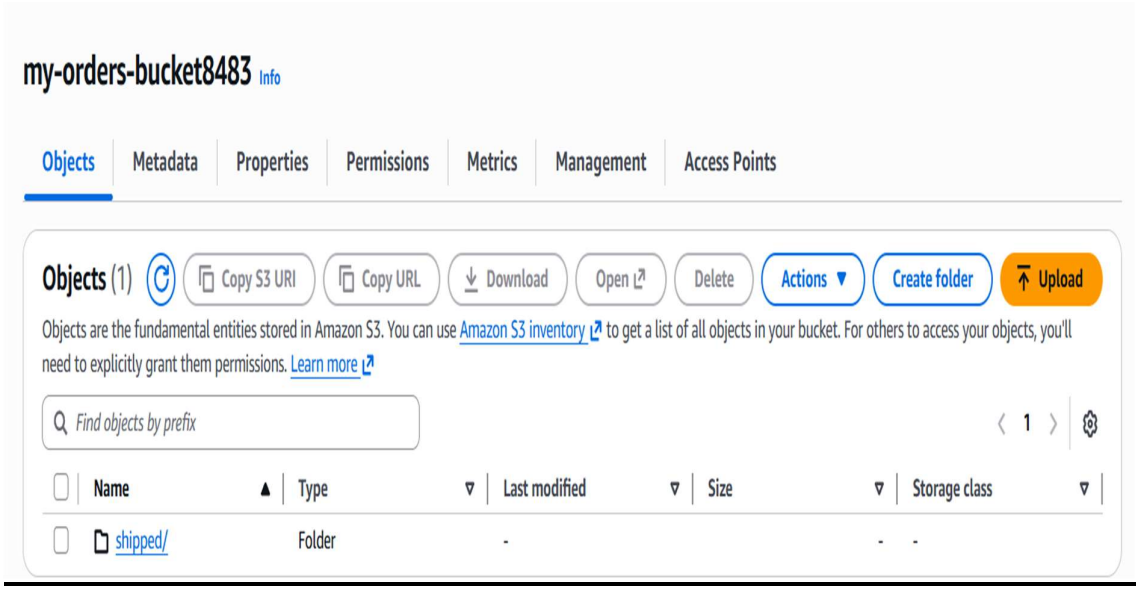

- Under shipped folder list of four items can show

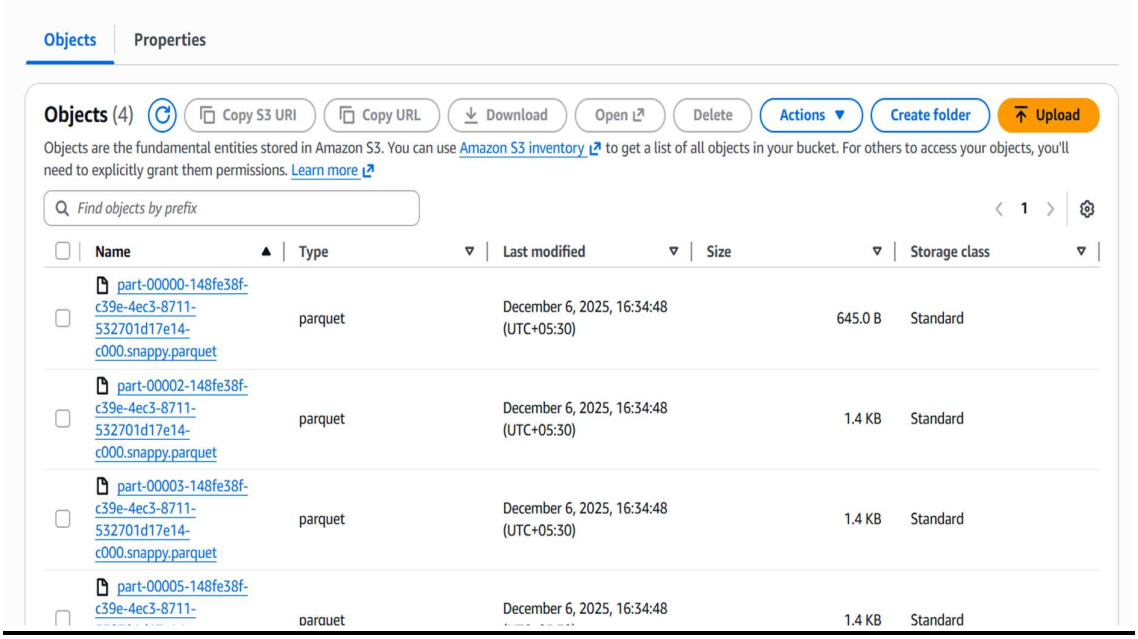

Step 8 : Athena Setup to run query

- AWS → Athena
- Select : Query your data in Athena console
- Launch Query editor
- Click on setting → edit your S3 location
- **s3://my-orders-bucket123/athena-results/** → Save
- perform query to create stored and view

1) Query :

```sql
SHOW DATABASES;
```

Its shows : **ordersdb**

- left side Database → select **ordersdb** into dropdown

2) Run this query to create table :

```sql
CREATE EXTERNAL TABLE IF NOT EXISTS ordersdb.orders_shipped (
  OrderID string,
  OrderDate string,
  Customer string,
  Amount bigint,
  Status string
)
STORED AS PARQUET
LOCATION 's3://my-orders-bucket8483/shipped/';
```

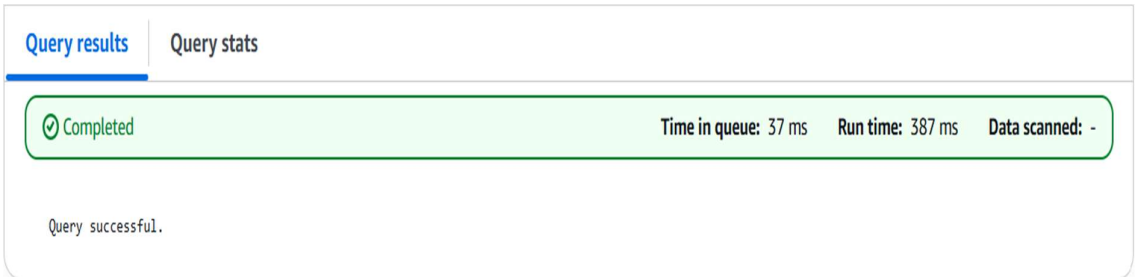

3) Query To view or fetch data :

```sql
SELECT * FROM ordersdb.orders_shipped
LIMIT 50;
```

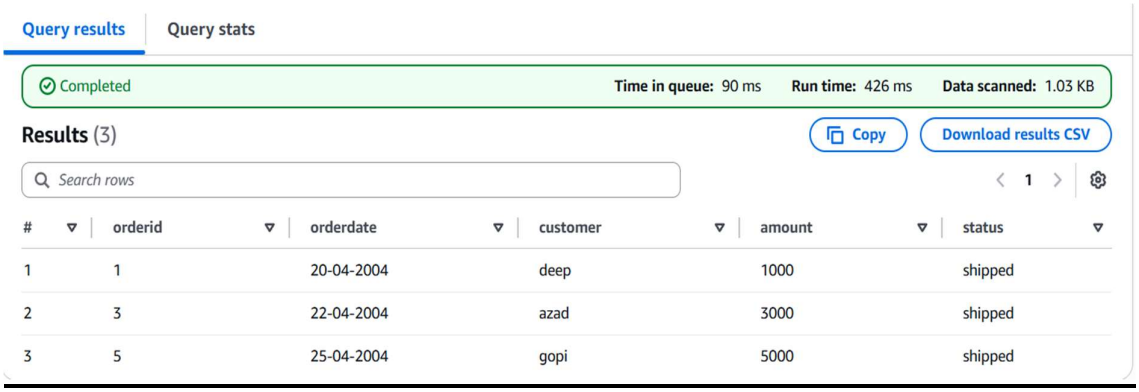

Query results | Query stats

Completed — Time in queue: 90 ms — Run time: 426 ms — Data scanned: 1.03 KB

Results (3) — Copy — Download results CSV

Search rows — 1

#	orderid	orderdate	customer	amount	status
1	1	20-04-2004	deep	1000	shipped
2	3	22-04-2004	azad	3000	shipped
3	5	25-04-2004	gopi	5000	shipped

4) Query To rows count

```sql
SELECT COUNT(*) AS total_orders FROM ordersdb.orders_shipped;
```

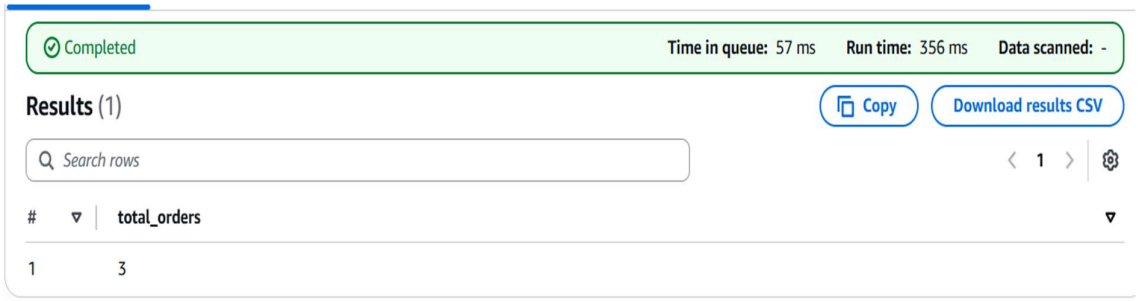

Completed — Time in queue: 57 ms — Run time: 356 ms — Data scanned: -

Results (1) — Copy — Download results CSV

Search rows — 1

#	total_orders
1	3

5) Query to Total amount (sum) overall

```sql
SELECT SUM(Amount) AS total_revenue FROM ordersdb.orders_shipped;
```

6) Revenue by customer (top customers)

```sql
SELECT Customer, SUM(Amount) AS total_spent
FROM ordersdb.orders_shipped
GROUP BY Customer
ORDER BY total_spent DESC
LIMIT 20;
```

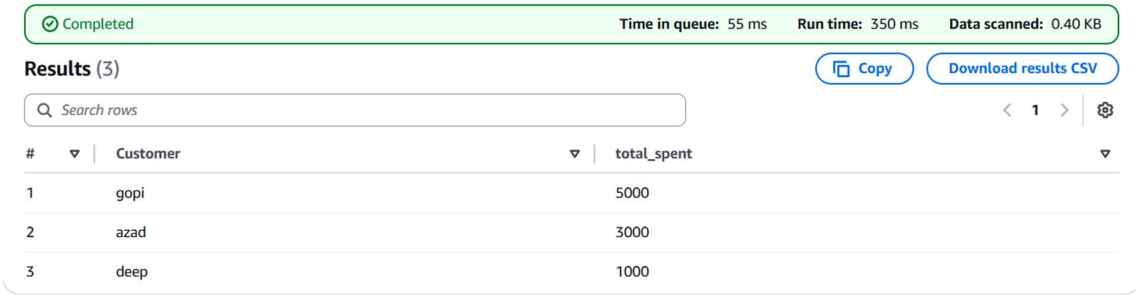

7) Orders per day (if OrderDate is YYYY-MM-DD)

```sql
SELECT OrderDate, COUNT(*) AS orders_count, SUM(Amount) AS daily_revenue
FROM ordersdb.orders_shipped
GROUP BY OrderDate
ORDER BY OrderDate;
```

Completed | Time in queue: 94 ms | Run time: 389 ms | Data scanned: 0.49 KB

Results (3) | Copy | Download results CSV

Search rows | 1

#	OrderDate	orders_count	daily_revenue
1	20-04-2004	1	1000
2	22-04-2004	1	3000
3	25-04-2004	1	5000

Final output :

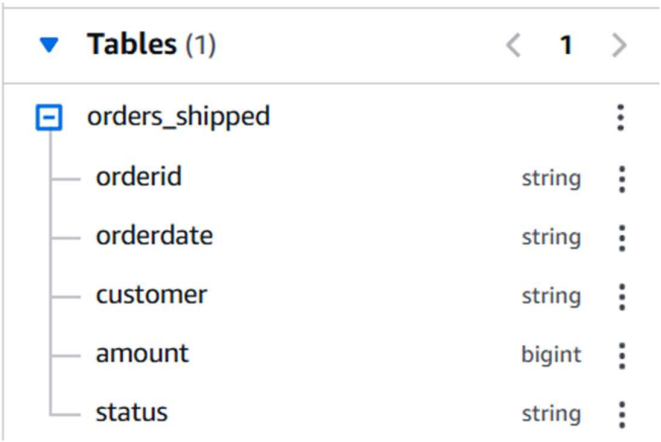

- **Purpose :**

  The purpose of this project is to build a fully automated and serverless data processing pipeline using AWS DynamoDB, AWS Glue, Amazon S3, and Amazon Athena. The aim is to extract raw order data from DynamoDB, automatically detect and catalog its schema using a Glue Crawler, transform and clean the data through a Glue ETL job, and store the optimized output in S3 using the Parquet format.

- **Conclusion :**

In conclusion, this project successfully demonstrates how AWS Glue and Amazon Athena can work together to create an end-to-end serverless data analytics solution. The pipeline extracted data from DynamoDB, transformed it using a Glue ETL job, stored it in an optimized Parquet format in S3, and allowed quick querying using Athena.

❖ Architecture Diagram :

Serverless Data Pipeline (DynamoDB , Glue, S3 , Athena)